

3/8/18

811

Distributed OS - distributed load (space distribution)

IPC

(1). Inter process communication.

(2). Same machine

RPC

(1). Remote Procedure call
(remote)

(2). Diff. mach.

Intraprocess communication - within a particular process (many threads) Multiple threads running within the process.

Obj.'l → Java is abide by fully OOP.
→ POP vs OOP.

4/8/18

Obj. Based → No polymorphism; Already classes are defined, Eg. - VR, ADN.

IPC: Mechanism that allows exchange of data b/w processes, by providing a user with a set of programming interfaces. IPC tells a programmer to organize the activities among diff. processes. It allows 1 applic. to control another thereby enabling data sharing without interference.

IPC mechanism can be classified into pipes (filtering), FIFO & shared memory.

" were introduced in UNIX OS where data flow is unidirectional. Working principle of FIFO is based on data space, FIFO FCFS basis. " also unidir. & is identified by access pts.

In shared mem. data sharing is done based on connection oriented or connection-less protocols.

In IPC data communication is allowed using semaphores, shared memory & other methods for info sharing. IPC facilitates efficient msg transfer b/w processes. The idea of IPC is based on task control architec (TCA). It is a flexible technique that can send & receive variable length arrays, data structures & lists.

Main IPC Methods:

Method

Description

File

A record stored on disk that can be accessed by name for any process.

Signal

A sys. msg. sent from 1 process to another, not usually stored but instead provide commands.

Socket

A data stream sent over a N/W interface either to a diff. process within same comp. or to another ".

Msg Queue

An anonymous data stream similar to a pipe, but stores & retrieves info in packets.

Pipe

A 2-way data stream interfaced thru std. I/O & O/P & is read char by char.

Named Pipe A pipe implemented thru a file on a file sys. instead of std. sys & C/F.

Semaphore A simple structure that synchronizes threads or processes acting on the shared resources.

Shared Mem. Multiple process given access to same memory allowing all to change & read the changes made by other process.

Msg-passing
(Shared nothing) Similar to msg, queue.

Mem. Mapped File A file mapped to the RAM & it can be modified by changing the memory address directly instead of O/Ping to a string. Stream. It shares the same benefits as a std. file.

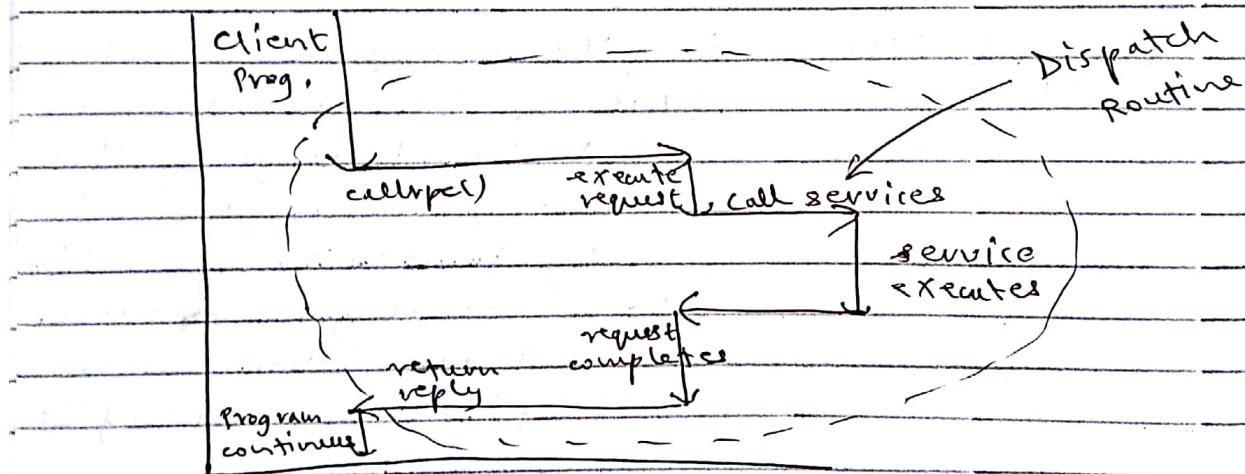
RPC: is also an IPC technology that allows a comp. prog. to call a subroutine or procedure to execute in another addr. space.
without the programmer explicitly coding the details for this remote interaction.

So RPC is just one kind of IPC viz a set of methods for exchange of data among multiple threads in 1/more processes.

The processes may be running in 1 or more comp. connec. by a Net but the IPC methods are divided into methods for msg. passing, synchronous, shared mem. & RPCs. The method of IPC here may vary based on the bandwidth & latency of communication b/w the threads & type of data being communicated. Several reasons for providing an envt. allowing process co-oper. are:

- ① info. sharing
- ② Speed up
- ③ Modularity
- ④ Convenience
- ⑤ Privilege separation

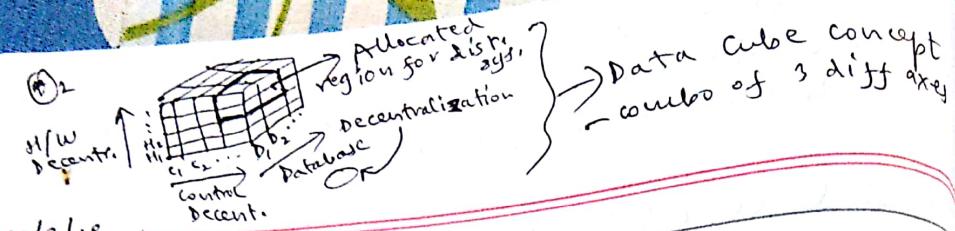
Q. How RPC works?



An RPC is analogous to a func. call like a func. call, when an RPC is met, the calling args. are passed to the RP & the caller waits for a response to be returned from the RP. The thread is blocked ^{from} processing until either a ~~reply~~[^] response is received or it times out.

When the request arrives, the server calls a dispatch routine that performs the requested service & sends the reply to the client. After the RPC call is completed the client prog. continues. RPC specifically supports N/W architec.

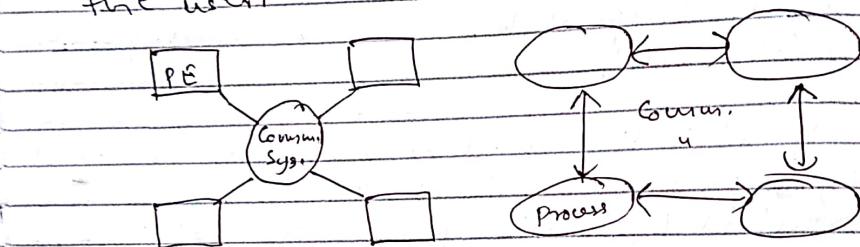
In the fig. the RPC mechanism is uniquely identified by a triplet (prog. no., version no. & procedure no.). " identifies group of related remote procedures which are unique, each prog. consists of multiple versions. " Version " is a collec. of procedures available to be called remotely.



SM

→ A distributed sys. is a collec. of independent computers that appear to the users of the sys. as a single comp.

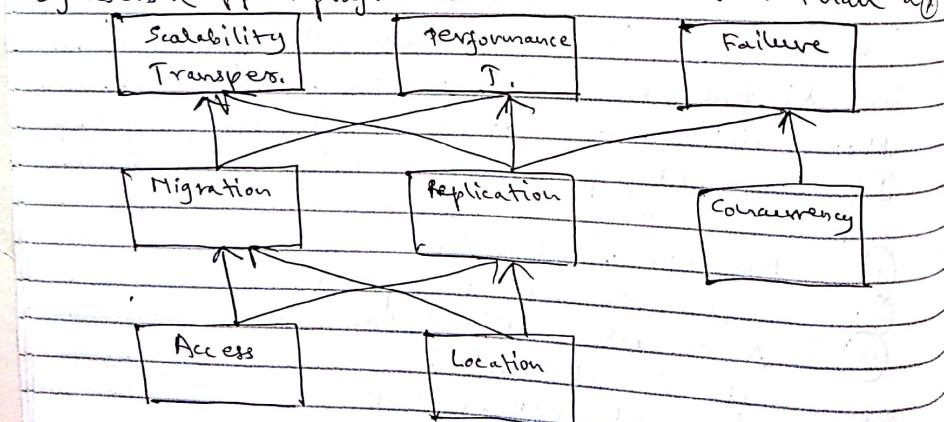
→ distributed systems are "seamless", the interfaces among functional units on the N/W are for the most part invisible to the user.



Fys. Sys.

Structure of → from physical view

② Transparency - distrib. sys. should be perceived by users & applic. programmers as a whole rather than as



↑ determines a bottom-up approach, i.e., the dependency

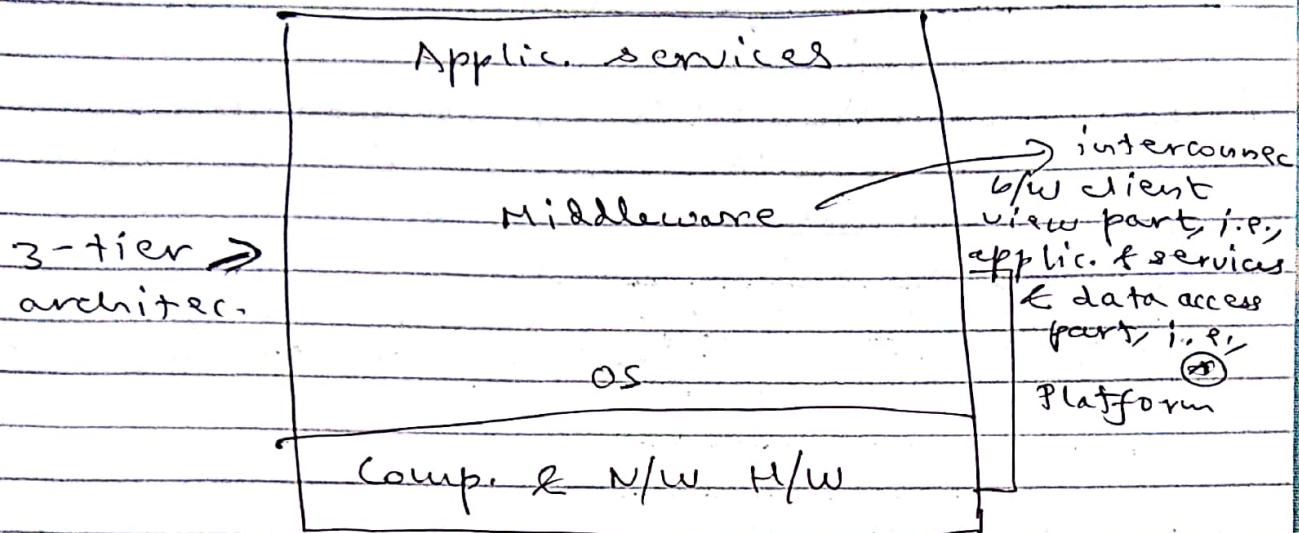
③ collec. of cooperating components. Transp. has diff. dimensions that were identified by ANSA. These rep. various p's

④ db part. It deals with business logic.

Thin server fat client - we are providing middleware along with client.

Fat & thin - student site, stud./customers are not allowed to modify any news or notice.

S/W & H/W service layers in distributed sys.



A working model of a distrib. sys. → ⑥ #

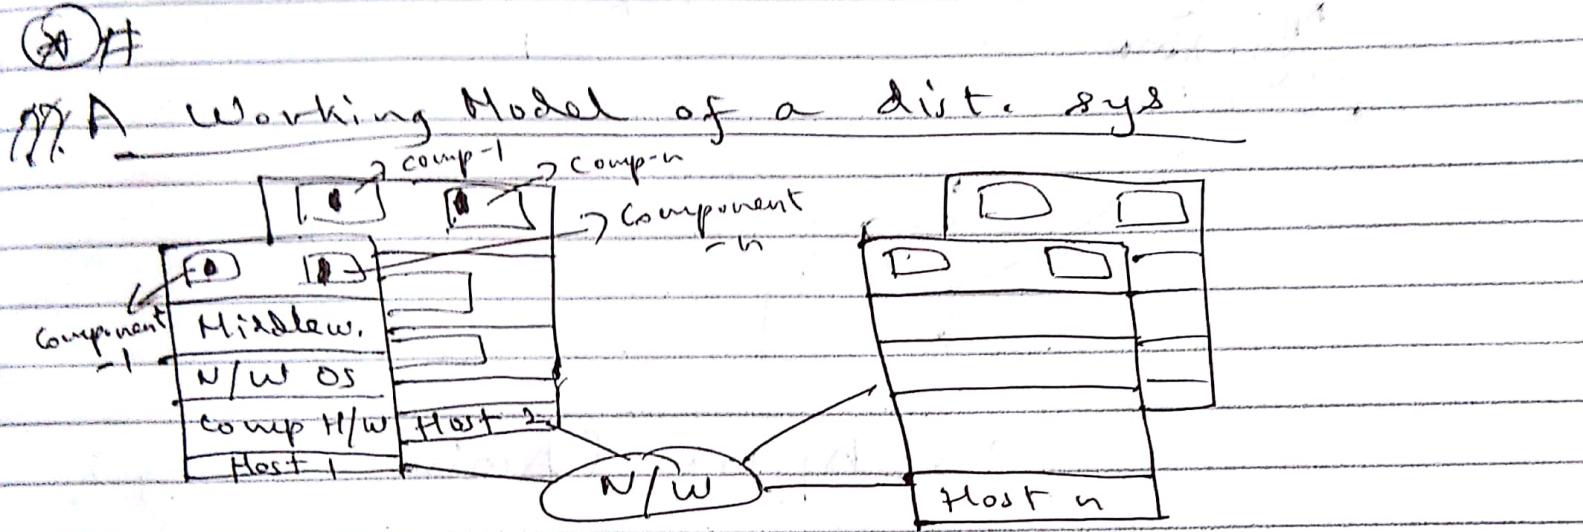
" service provided by multiple servers

" distributed applic. based on peer processes

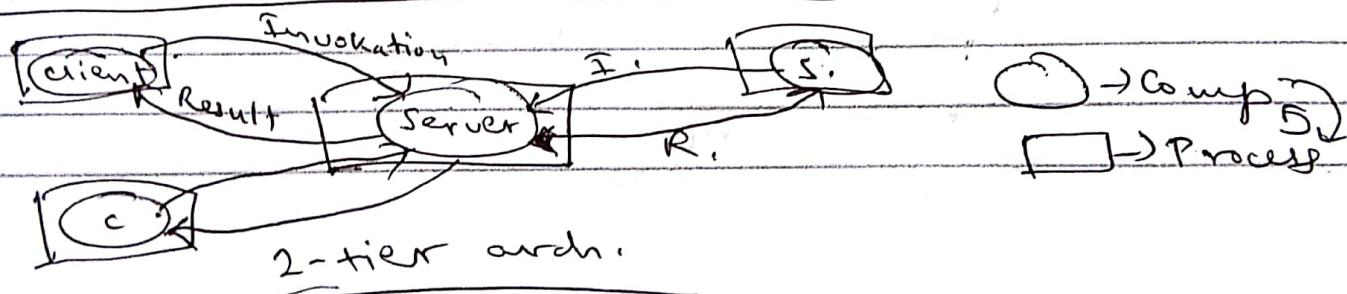
Web applets → active web pg.

The clients Thin clients & compute servers,

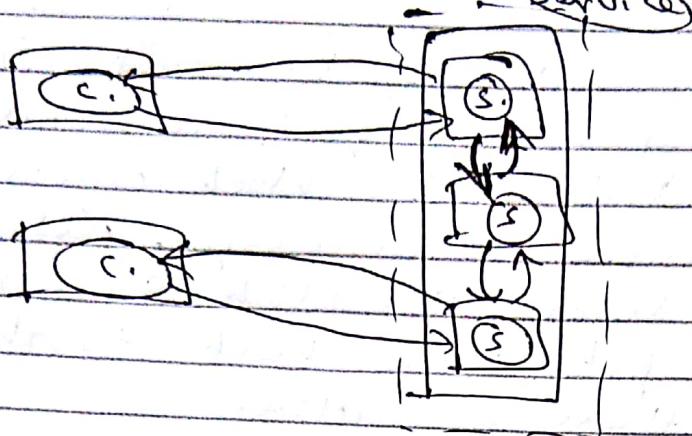
CIA upto fault tolerance



Clients invoke individ. servers

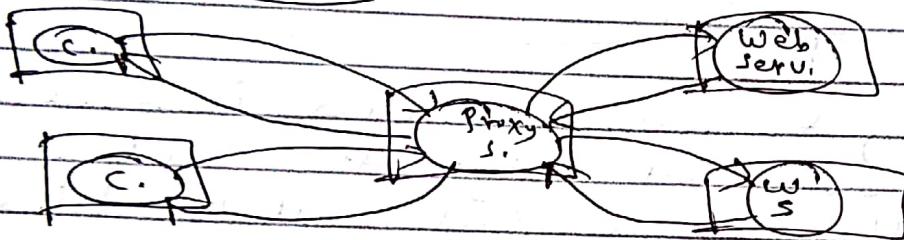


Proxy server
Backup " - When sys crashes

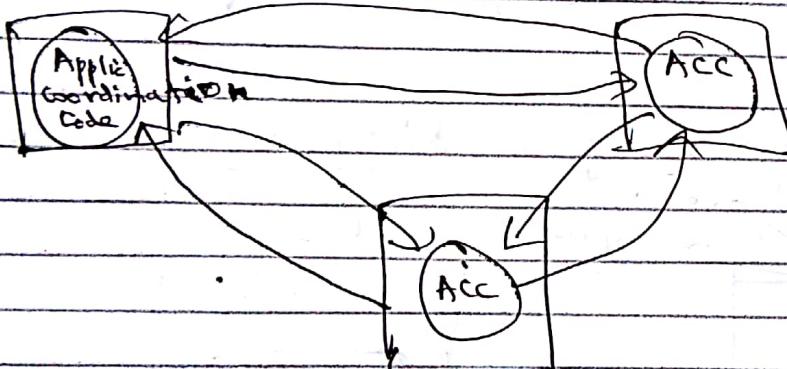


A service provided by multiple servers.

Web proxy s.



A dist. applic based on peer processes

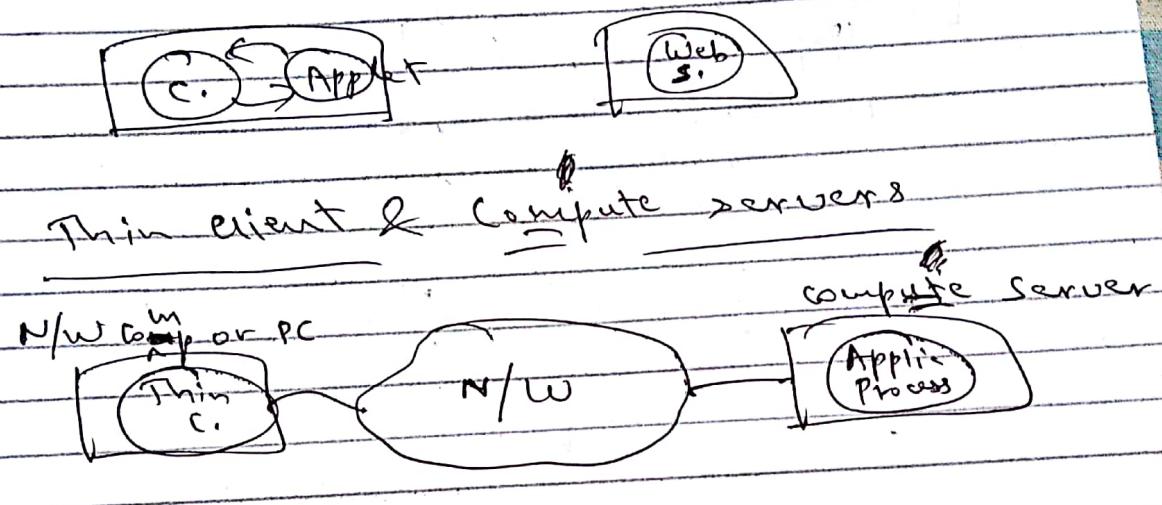


Web, Applets (Active web pg)

- a). Client request results in downloading of applet code



b) Client interacts with the applet



11/8/18

Fault Tolerance :

A system's ability to tolerate failure \rightarrow

1.

- > Reliability : likelihood that a sys. will remain operational for the duration of a mission.
 - requirement might be stated as 0.999999 availability for a 10 hr mission \rightarrow probab of failure during the mission must be at most 10^{-6} .
 - very high relab. is most imp in critical applic. - space shuttle, industrial controls in which failure could mean loss of life.

Sys. ability - 2
Availability; expresses the fraction of time
a sys. is operational.

→ A sys. with high availability may
in fact fail → its recovery time & failure
frequency must be small enough to
achieve the desired availab.

- high availab. is imp in airline reservations, telephone swit-
ching, etc, in which every minute of downtime translates into revenue
Importance of Design

→ A good fault-tolerant sys. design req.
a careful study of failures causes of
failures, & sys. responses to failures.
Such a study should be carried out in detail by the
design begins & must remain
part of the designing process.

Requirement Specification - I

4 - 2 ④

SM's notes (contd.):

④ Req. Specific.:

- Planning to avoid failures is most imp.
- A designer must analyze the env & determine the failures that must be tolerated to achieve the desired level of reliab.
- To optimise fault toler., it's imp to estimate actual failure rate for each possible failure.

Notee of SM

10/8/18:

Structure of a dist. sys.
↳ electronic struc.

	Other services		Security	
	file server			
Naming	Resource Alloc.	Deadlock Detect. Proc. Synchron.	res-our	comm. secur.
User names			-ces	
Name resolution	Proc. Mgt.			secu.
Address Routes	IPC			
	Transport protocol of supporting IPC primitives			

Classification of faults:

- > Transient " - occur once then disappear.
- > Intermittent " - " , goes away, then goes back,
- > Permanent faults - doesn't go away by itself, like disk failures.

Failure types:

- > Some are more probable than others
- > " transient others permanent
- > " occur in s/w, " in i/s/w,

Design issues of Fault Tolerant Sys.

Design - I

(1)

> Basic Principle - redundancy

— Spatial — redundant H/W

— Informational — " data struc.

— Temporal — " computation.

(2) Design of sys. that tolerate faults that occur while sys. is in use.

D - 2

> Redundancy costs money & time.

↳ sys. — one must optimize the design by trading off amount of redundancy used against level of fault tol.

SN's Notes:

Design Issues:

D-2 (contd.)

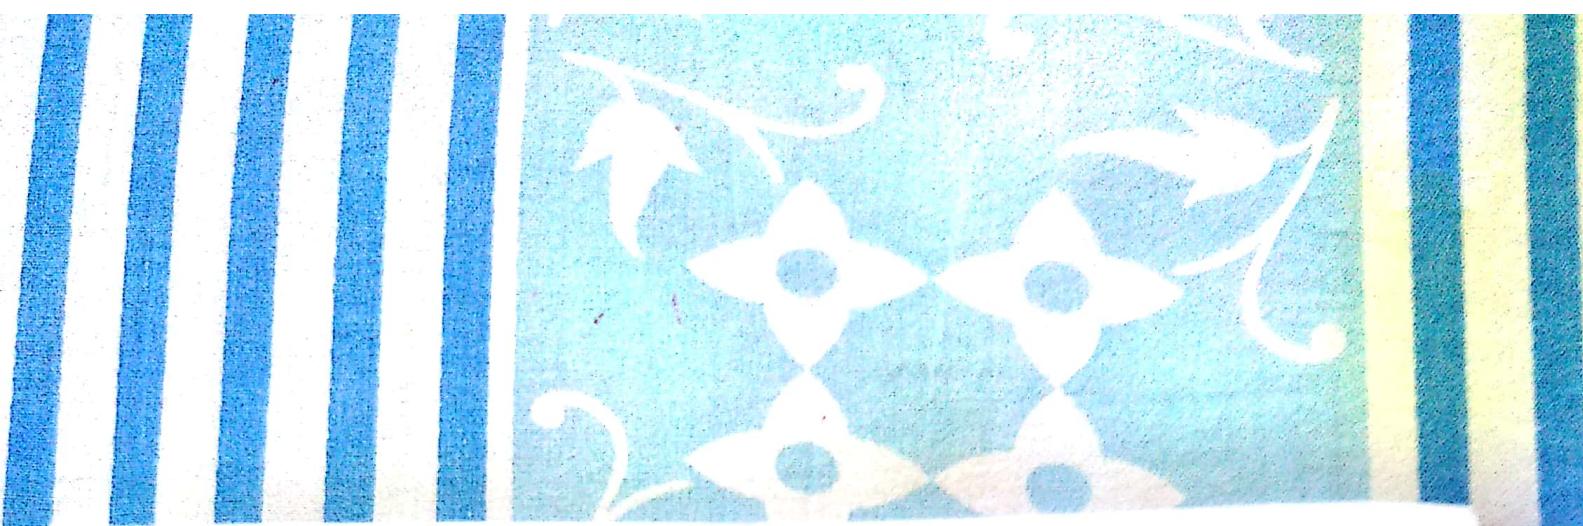
- Temporal redundancy usually requires more computation & it results in slower recovery from failure.
- Spatial has faster recovery but ↑ h/w costs, space, power, etc. requirements.

D-3

→ commonly used techniques for redundancy
— Modular redundancy.

D-3 (c.)

- Modular redundancy
 - Use multiple, identical replicas of H/W modules & a voter mechanism.
 - The o/p from the replicas are compared & correct o/p is determined - majority vote



* can tolerate most H/W faults that can affect minority of the H/W modules.

$$D = 4$$

→ n-version Programming $\textcircled{1}$ $\textcircled{2}$ $\textcircled{3}$

① ② ③

D - 4 (c.)

N - U. Prog.

- Write multiple versions of a s/w module
- O/Ps from these ⁺ are received & correct O/P is determined ^{via} voting mech.
- Each ver. is written by diff team, with the hope that they will not contain the same bugs
- Can tolerate s/w bugs that affect a minority of versions.

'can't tolerate' correlated fault - reason
for failure is common to 2 (or more)
modules eg. 2 modules share a single
power supply failure of which causes
both to fail.

Commit in SQL creates a checkpt.

Rollback - move back to the 4.

Error - Control Coding - 2

recovery blocks

18/8/18

SM

Dependability Evaluation - 1 :

> Once a sys. has been designed, it must be evaluated to determine if it meets reliability & dependability objectives.

> 2 " approaches:

> Use an Analytical Model

- can help developers to determine a sys. - ten's possible states & possibilities of transitions among them
- can be difficult to analyze models accurately

> Injecting faults

- Various types of faults can be injected to determine various dependability metrics.

Smoke testing - how many process can run at same time, Load testing, Others, crash/hang

Dependability Evaluation - 2 :

> In distributed systems a transac. based service can accept occasional failures followed by a lengthy recovery procedure

> A real-time service - Process control

- may have I/Os that are readings taken from sensors
- may have I/Os to actuators used to control a process directly or " activate alarm so that humans can intervene in the process.
- due to strict timing requirements recovery must be achieved within a v. small time limit, e.g., air traffic control, monitoring patients, controlling reactors,

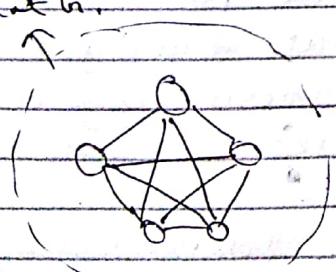
v.v. imp.

Fault Models

- > Omission Failure
 - A server omits to respond to a request or receive request
- > Response Failure
 - Value " - returns wrong val.
 - state transition failure - has wrong effect on resources
- > Timing Failure - any response that is not available to a client within a specified real time interval
- > Server crash failure - a server repeatedly fails to respond to requests until it's restarted
 - Amnesia - crash - a server starts in its initial state, having forgotten its state at the time of the crash, i.e., loses the values of the data items.
 - Pause - crash - a server restarts in the state before the crash.
 - Halting - crash - server never restarts.

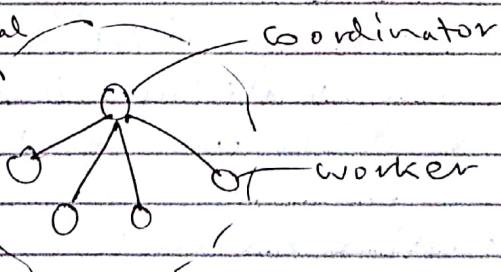
Flat Groups vs Hierarchical Gr.

flat gr.



communic. in a
flat group

Hierarchical
gr.



communic. in a
simple hierarchical
grp.

18

→ Agree. prob.

Agreement issue in faulty systems :

Possible assumptions about the underlying sys. :

1. Synchron. v. asynch. sys.
2. communic. delay is bounded or not
3. Msg. delivery is ordered or not
4. " transmission is done thru' unicasting or multicasting,

Circumst. under which distributed agreement can be reached if msg. ordering

^{unordered} _{ordered}

		Msg. Ordering				Bounded
		Un-	U.	B.	F.	
Syncrh.	A - u	X	X	X	X	Un- ord.
	U.			U.	M.	U.
Unicast Multic.						

Msg. transmission

Byzantine Agreement prob.:

Lamport et al. (1982)

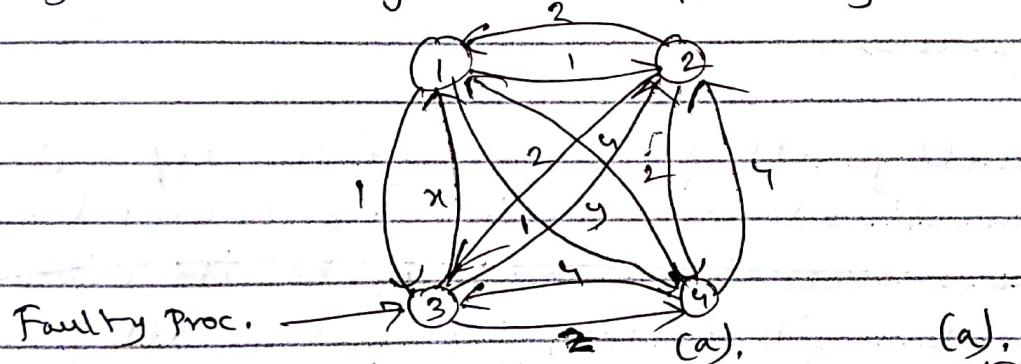
Assume reliable syncrh. ordered unicasted msg. sys. There are N processes, K of which may act as faulty or even malicious. A faulty process may send diff. val. to diff. processes.

Byzantine Failure: In fault-tolerant distributed computing, a B.T.F. is an arbitrary fault that occurs during exec. of an alg. in dist. sys. When a Byz. failure has occurred sys. may respond in an unpredictable way.

→ These arbitrary failures may be loosely categorized as :

- a failure to take another step in the algo, also known as a crash failure;
- a failure to correctly exec. a ~~the~~ step of "algo"; &
- arbitrary exec. of a step other than the one indicated by the algo.

Eg. Byzantine Agreement prob. for 4 processes



→ 3 non-faulty & 1 faulty process. Each proc. sends their val. to the others. Proc. 3 lies, giving diff val. x, y, z to diff proc.

24/8/18

SM

Byzantine Agreement Prob. - 2

1 Got	(1, 2, n, 4)	1 Got	2 Got	4 Got
2	" (1, 2, y, 4)	(1, 2, y, 4)	(1, 2, x, 4)	(1, 2, n, 4)
3	" (1, 2, 3, 4)	(a, b, c, d) ↳	(e, f, g, h)	(1, 2, y, 4)
4	" (1, 2, z, 4)	(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)
	(b)		(c)	

8. Replication & Checksum

(b) The vectors that each process assembles based on (a).

(c) The vec. " " " receives in step 3. Proc. 3 sends diff vectors to diff processes.

~~Vijay.~~
Byzantine Agreement prob - 3

> 3 processes can agree on the values received from 1, 2, 4. So that malicious proc. 3 val. is irrelevant..

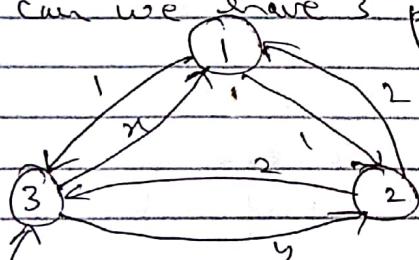
> If $N=3$ & $K=1$, i.e., only 2 non-faulty process - es this will not work.

> Lamport provided that, with $2K+1$ non-faulty proc., the sys. will survive K faulty proc., which makes a total of $3K+1$ proc. [Lamport, Alyo.]

BAP with 2 correct on faulty processes..

~~Q.~~ can we have 3 proc. having 1 faulty p.?

~~Tijp.~~



Faulty Proc.

(a)

1 not	(1, 2, n)
2 "	(1, 2, y)
3 "	(1, 2, 3)

(b)

1 not	2 not
(1, 2, y)	(1, 2, n)
(a, b, c)	(d, e, f)

(c)

> Fig. 8-6. The same as Fig. 8-5, except now with 2 correct processes & 1 faulty proc.

client server commun.

→ 2-tier arch. → if client
server commun.
course in exam

> TCP - Transport control protocol as a
connection oriented end-to-end commun.
unic. protocol is a reliable protocol.
But it does not prevent connection
establishment failure, which require searching
for new connections...

RPC semantics in the presence of failures:

> 5 diff classes of failures that can occur
in RPC systems:

1. The client is unable to locate the
server.

2. The request msg from the client to
the server is lost.

3. The server crashes after receiving
a request. This can be handled with
principles such as :

- At least once

- " most "

- The preferred principle is exactly
once, which is virtually impossible to
implement.

4. The reply msg from the server to
the client is lost.

5. The " crashes after sending a
request.

> Each case needs to be reserved prop-
erly to mask the failures.

Failure Egs.

➢ UDP service

- has omission failures coz it occasionally loses msgs.
- does not have val. failures coz it doesn't transmit corrupt msgs.

UPP uses checksum to mask the val. failures of the underlying IP by converting them to omission failures.

Reliable Multicasting

➢ Use -ve acknowledgement, known as scalable reliable multicasting - SRM

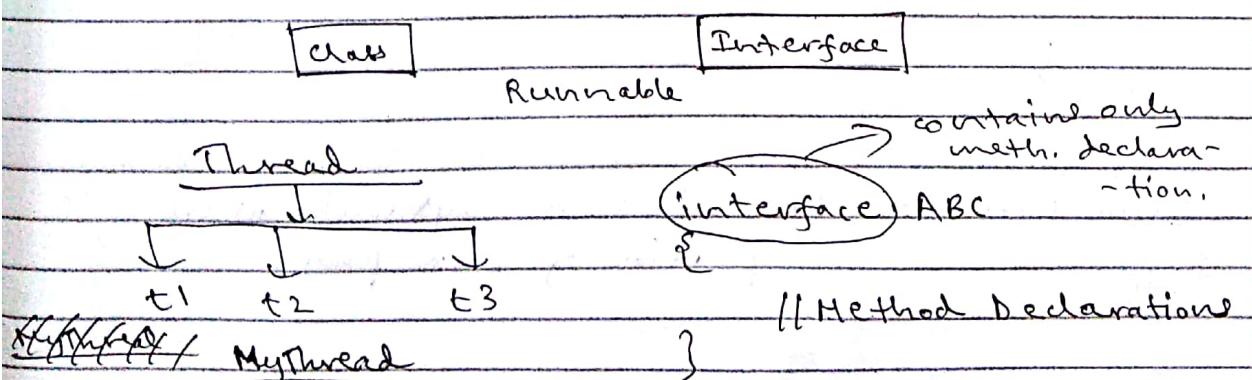
➢ Non-hierarchical & hierar. solns. are possible.

➢ Atomic multicasting requires all the replicas reaching agreement on the success or failure of multicast. This is known as dist. commit; two-phase or three-phase commit protocols can be used.

25/8/18

SM

➢ V. Sync. → Blocking, non-blocking
Threads: → intra process communic.



(2)

There's a class called Abstract class in Java.
Abs. Class is -

class A

{

 void getData()

{

 ;

}

 abstract void calculate();

}

ABC.java

+

→ Abc.java

class A

{

Abc.class

A.class

BC.u

 abstract void calculate();

}

interface BC

{

 abstract void c1();

 u

 u

 c2();

}

class Abc extends A implements BC

{

 public void psum(u)

{

}

Java becomes → class files or byte code after compilation of Java prog.

(3)

class MyThread extends Thread

{

 ps run(...)

 MyThread t1 = new MyThread();

 " "

 t2 = " "

better to write
implements Runnable

-able

 t1.start();

 t2.start();

 System.out.println("In thread 1 started");

 " ("In " " 2 " " ");

}

 p. v. run()

}

;

}

 and after all the code is written
 then we have to add the final part

 } final part

MIN_PRIORITY - 1 } }

NORM - 5 } }

MAX - 10 } }

→ Priorities

31/8/18

Thread :

Class → Thread

Interface → Runnable

```
int pl;  
pl = t1.getPriority();  
pl = pl + 2;  
t2.setPriority(pl);
```

class myThread extends Thread{

public void run(){

if(Thread.currentThread().isDaemon()) { // che-
// cking for daemon thread

~~Thread~~

System.out.println("daemon thread work");

}

else{

System.out.println("user thread work");

}

}

public static void main(...){

MyThread t1 = new MyThr(); // creating thread

" " t2 = " " ();

" " t3 = " " ();

t1.setDaemon(true); // now t1 is a daemon
thread

t1.start(); // starting thread

t2.start(); t3.start();

}

}

```
class NewThread implements Runnable
```

```
{
```

```
    String name;
```

```
    Thread t;
```

```
    NewThread(String threadname)
```

```
{
```

```
        name = threadname;
```

```
        t = new Thread(this, name);
```

```
        System.out.println("New Thread :" + t);
```

```
        t.start();
```

```
}
```

```
    public void run()
```

```
{
```

```
    try
```

```
    {
```

```
        for (int i = 5; i > 0; i--)
```

```
    {
```

```
        System.out.println(name + ":" + i);
```

```
        Thread.sleep(1000);
```

```
}
```

```
    catch (InterruptedException e)
```

```
{
```

```
        System.out.println(name + " interrupted");
```

```
}
```

```
        System.out.println(name + " exiting");
```

```
}
```

```
class MultiThreadDemo
```

```
{
```

```
    public void main(String args[])
```

```
    {
```

```
        new NewThread("One"),
```

```
        "Two",
```

```
        "Three",
```

```
try
{
    Thread.sleep(1000);
}
catch(InterruptedException e)
{
    System.out.println("Main Thread Interrupted.");
}
System.out.println("Exiting");
```

public class ThreadEx

```
{ sum(String args[])
{
```

```
    Thread t = new Thread();
    System.out.println("Now we are printing line after 2 seconds...");
```

```
    try
    {
```

```
        for(int i=1; i<=10; i++)
    {
```

```
            System.out.println(i);
```

```
            t.sleep(2000);
        }
```

```
}
```

```
    catch(InterruptedException e)
    {
```

```
        System.out.println("Thread interrupted...");
```

```
        e.printStackTrace();
    }
```

```
}
```

```
}
```

Error-control Coding

- Replic. is expensive
- For certain applic. - RAM, buses, error correcting codes can be used.
 - Hamming or other codes.
- Checkpts. & rollbacks
 - " are copies of applications' state saved in some storage that's immune to the failures under consideration.
- A rollback restarts the exec. from a previously saved checkpt.
- When a failure occurs, applic. state is rolled back to the prev. checkpt & restarted from there.
- Can be used to recover from transient as well as permanent H/w failures.
- Can be used for uniprocessor & distributed applic.

Recovery Blocks (for s/w failures)

- uses multiple alternates to perform the same func.
- One module is primary while others are secondary.
- When primary completes exec., its outcome is checked by an acceptance test.
- If O/P is not accep., a ^{-dary} ~~second~~ module executes until an acceptable O/P is obtained or alternates are exhausted.
- This meth. can tolerate s/w failures coz alternates are usually implemented with diff approaches (s/w algos.).

1/10/18

SM

Real Time Dist. Sys.:

- > In general the correctness of a prog depends on the logical sequence of instructions, not when they are executed.
- > A real-time programs & sys. interact with the external world in a way that involves time.
- > When a stimulus appears, the sys must respond to it in a certain way by a certain deadline.
- > If it delivers the correct ans but after deadline, the sys is regarded as having failed.

> Here ^{WHEN} the aw is produced is as imp as which aw is prod.

How a real time sys works?

- > Real time applic. & sys are highly structured much more so than general purpose dist sys.
- > Typically, an external device generates a stimulus for the comp, which ^{must} perform some action by a deadline.
- > When req. work has been completed the sys becomes idle, until next stimulus arrives.

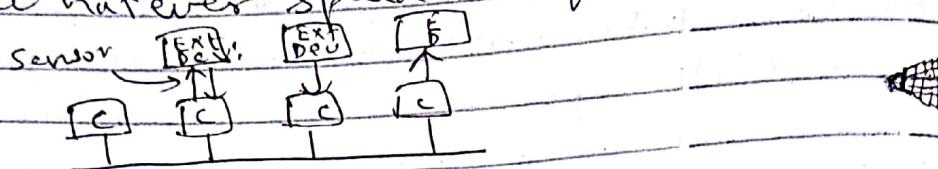
Types of stimuli:

- 1). periodic - with a stimulus occurring regularly every ΔT seconds such as a comp in a TV set or VCR getting a frame every $1/60$ of a sec.
- 2). Aperiodic - they are recurrent, but not regular as in the arrival of an aircraft in a air traffic controller's air space.
- 3). Sporadic (unexpect.) - such as a device overheating.

New Idea

- > There may happen that 2 events placed in the same time instant & to serve one ~~all~~ other will be neglec.

→ Some designers are experimenting with the idea of putting a dedicated microprocessor in front of each real time device to accept o/p from it whenever it has something to say, & give it i/p at whatever speed it requires.



Dist. Real time Sys.

→ A collec. of computers connected by a N/W.

→ Some of these are connected to ext. devices that produce or accept data or expect to be controlled in real time.

(a) → The comp. may be tiny micro controllers built into the devices or stand alone m/cs.

→ In both cases they usually have sensors for receiving signals from the devices &/or actuators for sending signals to them.

→ The sensors & actuators may be dig. or analog.

Soft real-time sys. - means that missing an occasional deadline is ~~not~~ alright. For eg, a telephone switch might be permitted to lose or misroute one call in T time per under Overload conditions & still be within specification.

Hard real-time sys. is unacceptable as this might lead to loss of life or an environmental catastrophe.

There are also informal sys. where missing a deadline means you have to kill off the current activity but the consequences is not fatal.

4/10/18

SM

Some Myths:

1. Real Time sys. are about writing device drivers in assembly code.
2. Real T. computing is fast computing.
- 3.

Design Issues:

1. Clock Synchronisation
2. Event-triggered vs Time-triggered sys.
3. Predictability
4. Fault Tolerance

Event-t-sys.:

- When a significant event in the outside world happens, it's detected by some sensor, which then cause the attached CPU to get an inter.

- 10
- It's an Inter. driven sys.
 - For soft real time system with lots of computing power to spare this appr. is simple, works well & is still widely used.
 - For complex sys, it is useful if comp can analyze the prog & know all there is to know about the sys behaviour once an event happens, even if it can't tell whether the event will happen.

Disadv: They can fail under conditions of heavy load that is when many events are happening at once.

(g) Time-trigg. sys :

→ Assume that all the clocks are synch. with a precision on the order of tens of ms. This prec. is possible as the protocol itself provides continuous clock sync. & has been designed to allow it to be in HW to extremely high prec.

- Ack inter. occurs every 2T milli sec.
- At each tick

Time t. d.:

- ΔT must be chosen with extreme care.
- If it's too small, the sys will get many clk inter & waste too much time finding them.
- If it's too large, errors may not be noticed until it's too late.
- Also the decision about which sensors to check on every clock tick & which to check on every other tick & so on is critical.
- Also, some events may be shorter than a clk tick so they must be saved to avoid losing them. They can be preserved electrically by latch ckts or by ^{embedded} microprocessor in the external devices.

Comparison:

- Event t. designs give faster response at low load but more overhead & chance of failure at high load.
- Time trig. designs have the opposite properties & are furthermore only suitable in a relatively static envt in which a great deal is known about sys behaviour in advance.

Predictability:

- real time system's behaviour is predictable.
- It should be clear at design time that the sys can meet all of its deadlines even at peak load.

Fault Tolerant: A sys that can
with operation without danger
is said to be fail-safe.

Lang - In which lang real time
sys will be interpreted \rightarrow study
from bk.

5/10/18

SM

TOKEN Ring with time limit

TDMA (Time Division Multiple Access):

A potential problem with wide
area real time distributed sys -
-tems is ~~their~~ relatively high
pkt loss rates. Std. protocol
deal with pkt loss by setting a
finer "when each" is transmitted. If
the "goes off" by the ack. is
received, the pkt is sent again.
One sol. to avoid it always to
transmit each pkt twice.

\hookrightarrow related to time-trig prot.

Real time scheduling:

$\begin{cases} \text{Dynamic} \\ \text{Static} \end{cases}$

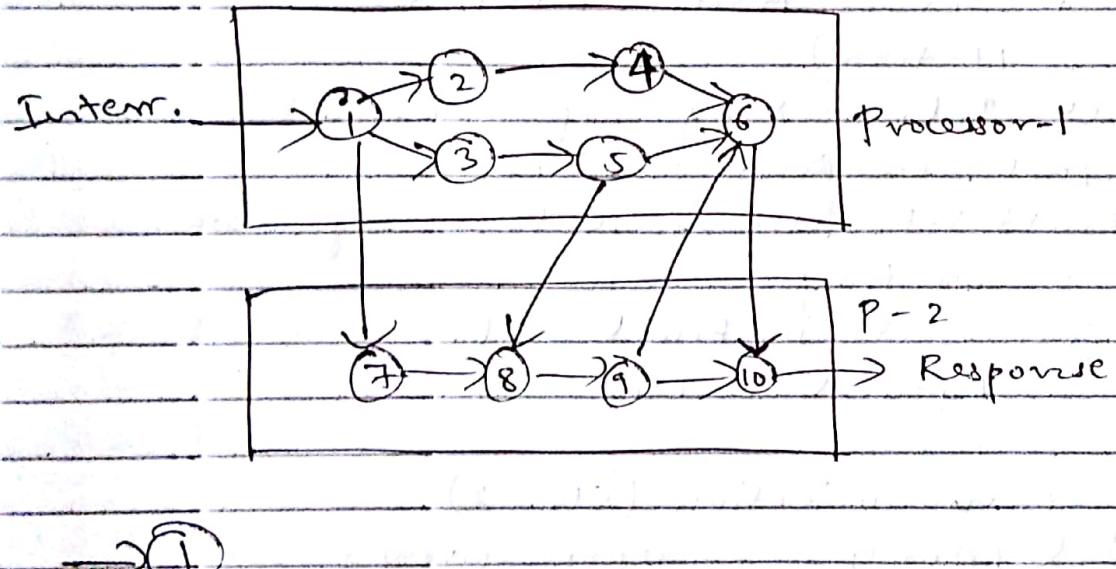
Dyn, S.: A periodic real time
dist. sys has n tasks & n
processors to run them on, let
the priority of the tasks be p_1, p_2, \dots, p_n

C_i be the CPU time needed by task i & let P_i is its period, i.e., the time b/w consecutive interv. To be feasible the utilization of the sys. U must be related to N as

$$U = \sum_{i=1}^n \frac{C_i}{P_i} \leq N$$

- types of scheduling for dist. dynamic scheduling real time s.:
- ① Rate Monotonic Alg.
 - ② Earliest Deadline First
 - ③ Laxity - How much CPU B_{T_i} pending for a process.

Static Sched.:



P_1	1	3	5	2	4	6
P_2	7		8	9		10